


# Towards Epistemic Machine Learning

## Uncertainty-aware ML

Uncertainty is inevitable and manifests itself at various stages of a learning problem. How should we incorporate that to build credible models?



## Explainable ML

How should we explain complex models? Can machine benefit from human explanations? What are the effects of explanations on human-AI collaborations?

## Preference Modelling

Learning from preference allows us to model and represent complex human behaviour. How should we incorporate that to promote model alignment?

